

CONTENT

PREFACE

xxiv

1 INTRODUCTION

1

1.1 WHAT IS AN OPERATING SYSTEM? 3

1.1.1 The Operating System as an Extended Machine 4

1.1.2 The Operating System as a Resource Manager 6

1.2 HISTORY OF OPERATING SYSTEMS 7

1.2.1 The First Generation (1945-55) Vacuum Tubes 7

1.2.2 The Second Generation (1955-65) Transistors and Batch Systems 8

1.2.3 The Third Generation (1965-1980) ICs and Multiprogramming 10

1.2.4 The Fourth Generation (1980-Present) Personal Computers 13

1.3 COMPUTER HARDWARE REVIEW 17

1.3.1 Processors 17

1.3.2 Memory 21

1.3.3 Disks 24

1.3.4 Tapes 25

1.3.5 I/O Devices 25

1.3.6 Buses 28

1.3.7 Booting the Computer 31

1.4 THE OPERATING SYSTEM ZOO	31
1.4.1 Mainframe Operating Systems	32
1.4.2 Server Operating Systems	32
1.4.3 Multiprocessor Operating Systems	32
1.4.4 Personal Computer Operating Systems	33
1.4.5 Handheld Computer Operating Systems	33
1.4.6 Embedded Operating Systems	33
1.4.7 Sensor Node Operating Systems	34
1.4.8 Real-Time Operating Systems	34
1.4.9 Smart Card Operating Systems	35
1.5 OPERATING SYSTEM CONCEPTS	35
1.5.1 Processes	36
1.5.2 Address Spaces	38
1.5.3 Files	38
1.5.4 Input/Output	41
1.5.5 Protection	42
1.5.6 The Shell	42
1.5.7 Ontogeny Recapitulates Phylogeny	44
1.6 SYSTEM CALLS	47
1.6.1 System Calls for Process Management	50
1.6.2 System Calls for File Management	54
1.6.3 System Calls for Directory Management	55
1.6.4 Miscellaneous System Calls	56
1.6.5 The Windows Win32 API	57
1.7 OPERATING SYSTEM STRUCTURE	60
1.7.1 Monolithic Systems	60
1.7.2 Layered Systems	61
1.7.3 Microkernels	62
1.7.4 Client-Server Model	65
1.7.5 Virtual Machines	65
1.7.6 Exokernels	69
1.8 THE WORLD ACCORDING TO C	70
1.8.1 The C Language	70
1.8.2 Header Files	71
1.8.3 Large Programming Projects	72
1.8.4 The Model of Run Time	73
1.9 RESEARCH ON OPERATING SYSTEMS	74

1.10 OUTLINE OF THE REST OF THIS BOOK	75
1.11 METRIC UNITS	76
1.12 SUMMARY	77

2 PROCESSES AND THREADS

2.1 PROCESSES	81
2.1.1 The Process Model	82
2.1.2 Process Creation	84
2.1.3 Process Termination	86
2.1.4 Process Hierarchies	87
2.1.5 Process States	88
2.1.6 Implementation of Processes	89
2.1.7 Modeling Multiprogramming	91
2.2 THREADS	93
2.2.1 Thread Usage	93
2.2.2 The Classical Thread Model	98
2.2.3 POSIX Threads	102
2.2.4 Implementing Threads in User Space	104
2.2.5 Implementing Threads in the Kernel	107
2.2.6 Hybrid Implementations	108
2.2.7 Scheduler Activations	109
2.2.8 Pop-Up Threads	110
2.2.9 Making Single-Threaded Code Multithreaded	112
2.3 INTERPROCESS COMMUNICATION	115
2.3.1 Race Conditions	115
2.3.2 Critical Regions	117
2.3.3 Mutual Exclusion with Busy Waiting	118
2.3.4 Sleep and Wakeup	123
2.3.5 Semaphores	126
2.3.6 Mutexes	128
2.3.7 Monitors	132
2.3.8 Message Passing	138
2.3.9 Barriers	142

2.4 SCHEDULING	143
2.4.1 Introduction to Scheduling	143
2.4.2 Scheduling in Batch Systems	150
2.4.3 Scheduling in Interactive Systems	152
2.4.4 Scheduling in Real-Time Systems	158
2.4.5 Policy versus Mechanism	159
2.4.6 Thread Scheduling	160
2.5 CLASSICAL IPC PROBLEMS	161
2.5.1 The Dining Philosophers Problem	162
2.5.2 The Readers and Writers Problem	165
2.6 RESEARCH ON PROCESSES AND THREADS	166
2.7 SUMMARY	167
 3 MEMORY MANAGEMENT	 173
3.1 NO MEMORY ABSTRACTION	174
3.2 A MEMORY ABSTRACTION: ADDRESS SPACES	177
3.2.1 The Notion of an Address Space	178
3.2.2 Swapping	179
3.2.3 Managing Free Memory	182
3.3 VIRTUAL MEMORY	186
3.3.1 Paging	187
3.3.2 Page Tables	191
3.3.3 Speeding Up Paging	192
3.3.4 Page Tables for Large Memories	196
3.4 PAGE REPLACEMENT ALGORITHMS	199
3.4.1 The Optimal Page Replacement Algorithm	200
3.4.2 The Not Recently Used Page Replacement Algorithm	201
3.4.3 The First-In, First-Out (FIFO) Page Replacement Algorithm	202
3.4.4 The Second-Chance Page Replacement Algorithm	202
3.4.5 The Clock Page Replacement Algorithm	203
3.4.6 The Least Recently Used (LRU) Page Replacement Algorithm	204
3.4.7 Simulating LRU in Software	205
3.4.8 The Working Set Page Replacement Algorithm	207

3.4.9 The WSClock Page Replacement Algorithm	211
3.4.10 Summary of Page Replacement Algorithms	213
3.5 DESIGN ISSUES FOR PAGING SYSTEMS	214
3.5.1 Local versus Global Allocation Policies	214
3.5.2 Load Control	216
3.5.3 Page Size	217
3.5.4 Separate Instruction and Data Spaces	219
3.5.5 Shared Pages	219
3.5.6 Shared Libraries	221
3.5.7 Mapped Files	223
3.5.8 Cleaning Policy	224
3.5.9 Virtual Memory Interface	224
3.6 IMPLEMENTATION ISSUES	225
3.6.1 Operating System Involvement with Paging	225
3.6.2 Page Fault Handling	226
3.6.3 Instruction Backup	227
3.6.4 Locking Pages in Memory	228
3.6.5 Backing Store	229
3.6.6 Separation of Policy and Mechanism	231
3.7 SEGMENTATION	232
3.7.1 Implementation of Pure Segmentation	235
3.7.2 Segmentation with Paging: MULTICS	236
3.7.3 Segmentation with Paging: The Intel Pentium	240
3.8 RESEARCH ON MEMORY MANAGEMENT	245
3.9 SUMMARY	246
 4 FILE SYSTEMS	 255
4.1 FILES	255
4.1.1 Fit- W-ming	255
4.1.2 Figgcture	257
4.1.3 Fii^^ws	258
4.1.4 File Access	260
4.1.5 File Attributes	261

4.1.6 File Operations	262
4.1.7 An Example Program Using File System Calls	263
4.2 DIRECTORIES	266
4.2.1 Single-Level Directory Systems	266
4.2.2 Hierarchical Directory Systems	266
4.2.3 Path Names	267
4.2.4 Directory Operations	270
4.3 FILE SYSTEM IMPLEMENTATION	271
4.3.1 File System Layout	271
4.3.2 Implementing Files	272
4.3.3 Implementing Directories	278
4.3.4 Shared Files	281
4.3.5 Log-Structured File Systems	283
4.3.6 Journaling File Systems	285
4.3.7 Virtual File Systems	286
4.4 FILE SYSTEM MANAGEMENT AND OPTIMIZATION	290
4.4.1 Disk Space Management	290
4.4.2 File System Backups	296
4.4.3 File System Consistency	302
4.4.4 File System Performance	305
4.4.5 Defragmenting Disks	309
4.5 EXAMPLE FILE SYSTEMS	310
4.5.1 CD-ROM File Systems	310
4.5.2 The MS-DOS File System	316
4.5.3 The UNIX V7 File System	319
4.6 RESEARCH ON FILE SYSTEMS	322
4.7 SUMMARY	322

5 INPUT/OUTPUT

5.1 PRINCIPLES OF I/O HARDWARE	327
5.1.1 I/O Devices	328
5.1.2 Device Controllers	329

5.1.3 Memory-Mapped I/O	330
5.1.4 Direct Memory Access (DMA)	334
5.1.5 Interrupts Revisited	337
5.2 PRINCIPLES OF I/O SOFTWARE	341
5.2.1 Goals of the I/O Software	341
5.2.2 Programmed I/O	342
5.2.3 Interrupt-Driven I/O	344
5.2.4 I/O Using DMA	345
5.3 I/O SOFTWARE LAYERS	346
5.3.1 Interrupt Handlers	346
5.3.2 Device Drivers	347
5.3.3 Device-Independent I/O Software	351
5.3.4 User-Space I/O Software	357
5.4 DISKS	358
5.4.1 Disk Hardware	359
5.4.2 Disk Formatting	374
5.4.3 Disk Arm Scheduling Algorithms	377
5.4.4 Error Handling	380
5.4.5 Stable Storage	383
5.5 CLOCKS	386
5.5.1 Clock Hardware	386
5.5.2 Clock Software	388
5.5.3 Soft Timers	391
5.6 USER INTERFACES: KEYBOARD, MOUSE, MONITOR	392
5.6.1 Input Software	392
5.6.2 Output Software	397
5.7 THIN CLIENTS	413
5.8 POWER MANAGEMENT	415
5.8.1 Hardware Issues	416
5.8.2 Operating System Issues	417
5.8.3 Application Program Issues	422
5.9 RESEARCH ON INPUT/OUTPUT	423
5.10 SUMMARY	424

6 DEADLOCKS

6.1 RESOURCES	432
6.1.1 Preemptable and Nonpreemptable Resources	432
6.1.2 Resource Acquisition	433
6.2 INTRODUCTION TO DEADLOCKS	435
6.2.1 Conditions for Resource Deadlocks	436
6.2.2 Deadlock Modeling	436
6.3 THE OSTRICH ALGORITHM	439
6.4 DEADLOCK DETECTION AND RECOVERY	440
6.4.1 Deadlock Detection with One Resource of Each Type	440
6.4.2 Deadlock Detection with Multiple Resources of Each Type	
6.4.3 Recovery from Deadlock	445
6.5 DEADLOCK AVOIDANCE	446
6.5.1 Resource Trajectories	447
6.5.2 Safe and Unsafe States	448
6.5.3 The Banker's Algorithm for a Single Resource	449
6.5.4 The Banker's Algorithm for Multiple Resources	450
6.6 DEADLOCK PREVENTION	452
6.6.1 Attacking the Mutual Exclusion Condition	452
6.6.2 Attacking the Hold and Wait Condition	453
6.6.3 Attacking the No Preemption Condition	453
6.6.4 Attacking the Circular Wait Condition	454
6.7 OTHER ISSUES	455
6.7.1 Two-Phase Locking	455
6.7.2 Communication Deadlocks	456
6.7.3 Livelock	457
6.7.4 Starvation	459
6.8 RESEARCH ON DEADLOCKS	459
6.9 SUMMARY	460

7 MULTIMEDIA OPERATING SYSTEMS

7.1 INTRODUCTION TO MULTIMEDIA	466
7.2 MULTIMEDIA FILES	470
7.2.1 Video Encoding	471
7.2.2 Audio Encoding	474
7.3 VIDEO COMPRESSION	476
7.3.1 The JPEG Standard	476
7.3.2 The MPEG Standard	479
7.4 AUDIO COMPRESSION	482
7.5 MULTIMEDIA PROCESS SCHEDULING	485
7.5.1 Scheduling Homogeneous Processes	486
7.5.2 General Real-Time Scheduling	486
7.5.3 Rate Monotonic Scheduling	488
7.5.4 Earliest Deadline First Scheduling	489
7.6 MULTIMEDIA FILE SYSTEM PARADIGMS	491
7.6.1 VCR Control Functions	492
7.6.2 Near Video on Demand	494
7.6.3 Near Video on Demand with VCR Functions	496
7.7 FILE PLACEMENT	497
7.7.1 Placing a File on a Single Disk	498
7.7.2 Two Alternative File Organization Strategies	499
7.7.3 Placing Files for Near Video on Demand	502
7.7.4 Placing Multiple Files on a Single Disk	504
7.7.5 Placing Files on Multiple Disks	506
7.8 CACHING	508
7.8.1 Block Caching	509
7.8.2 File Caching	510
7.9 DISK SCHEDULING FOR MULTIMEDIA	511
7.9.1 Static Disk Scheduling	511
7.9.2 Dynamic Disk Scheduling	513
7.10 RESEARCH ON MULTIMEDIA	514
7.11 SUMMARY	515

8 MULTIPLE PROCESSOR SYSTEMS

521

8.1 MULTIPROCESSORS	524
8.1.1 Multiprocessor Hardware	524
8.1.2 Multiprocessor Operating System Types	532
8.1.3 Multiprocessor Synchronization	536
8.1.4 Multiprocessor Scheduling	540
8.2 MULTICOMPUTERS	546
8.2.1 Multicomputer Hardware	547
8.2.2 Low-Level Communication Software	551
8.2.3 User-Level Communication Software	553
8.2.4 Remote Procedure Call	556
8.2.5 Distributed Shared Memory	558
8.2.6 Multicomputer Scheduling	563
8.2.7 Load Balancing	563
8.3 VIRTUALIZATION	566
8.3.1 Requirements for Virtualization	568
8.3.2 Type 1 Hypervisors	569
8.3.3 Type 2 Hypervisors	570
8.3.4 Paravirtualization	572
8.3.5 Memory Virtualization	574
8.3.6 I/O Virtualization	576
8.3.7 Virtual Appliances	577
8.3.8 Virtual Machines on Multicore CPUs	577
8.3.9 Licensing Issues	578
8.4 DISTRIBUTED SYSTEMS	578
8.4.1 Network Hardware	581
8.4.2 Network Services and Protocols	584
8.4.3 Document-Based Middleware	588
8.4.4 File-System-Based Middleware	589
8.4.5 Object-Based Middleware	594
8.4.6 Coordination-Based Middleware	596
8.4.7 Grids	601
8.5 RESEARCH ON MULTIPLE PROCESSOR SYSTEMS	602
8.6 SUMMARY	603

9 SECURITY

9.1 THE SECURITY ENVIRONMENT	611
9.1.1 Threats	611
9.1.2 Intruders	613
9.1.3 Accidental Data Loss	614
9.2 BASICS OF CRYPTOGRAPHY	614
9.2.1 Secret-Key Cryptography	615
9.2.2 Public-Key Cryptography	616
9.2.3 One-Way Functions	617
9.2.4 Digital Signatures	617
9.2.5 Trusted Platform Module	619
9.3 PROTECTION MECHANISMS	620
9.3.1 Protection Domains	620
9.3.2 Access Control Lists	622
9.3.3 Capabilities	625
9.3.4 Trusted Systems	628
9.3.5 Trusted Computing Base	629
9.3.6 Formal Models of Secure Systems	630
9.3.7 Multilevel Security	632
9.3.8 Covert Channels	635
9.4 AUTHENTICATION	639
9.4.1 Authentication Using Passwords	640
9.4.2 Authentication Using a Physical Object	649
9.4.3 Authentication Using Biometrics	651
9.5 INSIDER ATTACKS	654
9.5.1 Logic Bombs	654
9.5.2 Trap Doors	655
9.5.3 Login Spooling	656
9.6 EXPLOITING CODE BUGS	657
9.6.1 Buffer Overflow Attacks	658
9.6.2 Format String Attacks	660
9.6.3 Return to libc Attacks	662
9.6.4 Integer Overflow Attacks	663
9.6.5 Code Injection Attacks	664
9.6.6 Privilege Escalation Attacks	665

9.7 MALWARE	665
9.7.1 Trojan Horses	668
9.7.2 Viruses	670
9.7.3 Worms	680
9.7.4 Spyware	682
9.7.5 Rootkits	686
9.8 DEFENSES	690
9.8.1 Firewalls	691
9.8.2 Antivirus and Anti-Antivirus Techniques	693
9.8.3 Code Signing	699
9.8.4 Jailing	700
9.8.5 Model-Based Intrusion Detection	701
9.8.6 Encapsulating Mobile Code	703
9.8.7 Java Security	707
9.9 RESEARCH ON SECURITY	709
9.10 SUMMARY	710

10 CASE STUDY 1: LINUX

715

10.1 HISTORY OF UNIX AND LINUX	716
10.1.1 UNICS	716
10.1.2 PDP-11 UNIX	717
10.1.3 Portable UNIX	718
10.1.4 Berkeley UNIX	719
10.1.5 Standard UNIX	720
10.1.6 MINTX	721
10.1.7 Linux	722
10.2 OVERVIEW OF LINUX	724
10.2.1 Linux Goals	725
10.2.2 Interfaces to Linux	726
10.2.3 The Shell	727
10.2.5 Kernel Structure	732

10.3 PROCESSES IN LINUX	735
10.3.1 Fundamental Concepts	735
10.3.2 Process Management System Calls in Linux	737
10.3.3 Implementation of Processes and Threads in Linux	741
10.3.4 Scheduling in Linux	748
10.3.5 Booting Linux	751
10.4 MEMORY MANAGEMENT IN LINUX	754
10.4.1 Fundamental Concepts	754
10.4.2 Memory Management System Calls in Linux	757
10.4.3 Implementation of Memory Management in Linux	758
10.4.4 Paging in Linux	764
10.5 INPUT/OUTPUT IN LINUX	767
10.5.1 Fundamental Concepts	768
10.5.2 Networking	769
10.5.3 Input/Output System Calls in Linux	771
10.5.4 Implementation of Input/Output in Linux	771
10.5.5 Modules in Linux	775
10.6 THE LINUX FILE SYSTEM	775
10.6.1 Fundamental Concepts	776
10.6.2 File System Calls in Linux	781
10.6.3 Implementation of the Linux File System	784
10.6.4 NFS: The Network File System	792
10.7 SECURITY IN LINUX	799
10.7.1 Fundamental Concepts	799
10.7.2 Security System Calls in Linux	801
10.7.3 Implementation of Security in Linux	802
10.8 SUMMARY	802
11 CASE STUDY 2: WINDOWS VISTA	809
11.1 HISTORY OF WINDOWS VISTA	809
11.1.1 1980s: MS-DOS	810
11.1.2 1990s: MS-DOS-based Windows	811
11.1.3 2000s: NT-based Windows	811
11.1.4 Windows Vista	814

11.2 PROGRAMMING WINDOWS VISTA	815
11.2.1 The Native NT Application Programming Interface	818
11.2.2 The Win32 Application Programming Interface	821
11.2.3 The Windows Registry	825
11.3 SYSTEM STRUCTURE	827
11.3.1 Operating System Structure	828
11.3.2 Booting Windows Vista	843
11.3.3 Implementation of the Object Manager	844
11.3.4 Subsystems, DLLs, and User-Mode Services	854
11.4 PROCESSES AND THREADS IN WINDOWS VISTA	857
11.4.1 Fundamental Concepts	857
11.A2 Job, Process, Thread, and Fiber Management API Calls	862
11.4.3 Implementation of Processes and Threads	867
11.5 MEMORY MANAGEMENT	875
11.5.1 Fundamental Concepts	875
11.5.2 Memory Management System Calls	880
11.5.3 Implementation of Memory Management	881
11.6 CACHING IN WINDOWS VISTA	890
11.7 INPUT/OUTPUT IN WINDOWS VISTA	892
11.7.1 Fundamental Concepts	893
11.7.2 Input/Output API Calls	894
11.7.3 Implementation of I/O	897
11.8 THE WINDOWS NT FILE SYSTEM	902
11.8.1 Fundamental Concepts	903
11.8.2 Implementation of the NT File System	904
11.9 SECURITY IN WINDOWS VISTA	914
11.9.1 Fundamental Concepts	915
11.9.2 Security API Calls	917
11.9.3 Implementation of Security	918
11.10 SUMMARY	920

12 CASE STUDY 3: SYMBIAN OS

12.1 THE HISTORY OF SYMBIAN OS	926
12.1.1 Symbian OS Roots: Psion and EPOC	926
12.1.2 Symbian OS Version 6	927
12.1.3 Symbian OS Version 7	928
12.1A Symbian OS Today	928
12.2 AN OVERVIEW OF SYMBIAN OS	928
12.2.1 Object Orientation	929
12.2.2 Microkernel Design	930
12.2.3 The Symbian OS Nanokernel	931
12.2.4 Client/Server Resource Access	931
12.2.5 Features of a Larger Operating System	932
12.2.6 Communication and Multimedia	933
12.3 PROCESSES AND THREADS IN SYMBIAN OS	933
12.3.1 Threads and Nanothreads	934
12.3.2 Processes	935
12.3.3 Active Objects	935
12.3.4 Interprocess Communication	936
12.4 MEMORY MANAGEMENT	937
12.4.1 Systems with No Virtual Memory	937
12.4.2 How Symbian OS Addresses Memory	939
12.5 INPUT AND OUTPUT	941
12.5.1 Device Drivers	941
12.5.2 Kernel Extensions	942
12.5.3 Direct Memory Access	942
12.5.4 Special Case: Storage Media	943
12.5.5 Blocking I/O	943
12.5.6 Removable Media	944
12.6 STORAGE SYSTEMS	944
12.6.1 File Systems for Mobile Devices	944
12.6.2 Symbian OS File Systems	945
12.6.3 File System Security and Protection	945
12.7 SECURITY IN SYMBIAN OS	946

12.8 COMMUNICATION IN SYMBIAN OS 949

12.8.1 Basic Infrastructure 949

12.8.2 A Closer Look at the Infrastructure 950

12.9 SUMMARY 953

13 OPERATING SYSTEM DESIGN

13.1 THE NATURE OF THE DESIGN PROBLEM 956

13.1.1 Goals 956

13.1.2 Why Is It Hard to Design an Operating System? 957

13.2 INTERFACE DESIGN 959

13.2.1 Guiding Principles 959

13.2.2 Paradigms 961

13.2.3 The System Call Interface 964

13.3 IMPLEMENTATION 967

13.3.1 System Structure 967

13.3.2 Mechanism versus Policy 971

13.3.3 Orthogonality 972

13.3.4 Naming 973

13.3.5 Binding Time 974

13.3.6 Static versus Dynamic Structures 975

13.3.7 Top-Down versus Bottom-Up Implementation 976

13.3.8 Useful Techniques 977

13.4 PERFORMANCE 983

13.4.1 Why Are Operating Systems Slow? 983

13.4.2 What Should Be Optimized? 984

13.4.3 Space-Time Trade-offs 984

13.4.4 Caching 987

13.4.5 Hints 988

13.4.6 Exploiting Locality 989

13.4.7 Optimize the Common Case 989

13.5 PROJECT MANAGEMENT 990

13.5.1 The Mythical Man Month 990

13.5.2 Team Structure 991

13.5.3 The Role of Experience 993

13.5.4 No Silver Bullet 994

13.6 TRENDS IN OPERATING SYSTEM DESIGN 994

13.6.1 Virtualization 995

13.6.2 Multicore Chips 995

13.6.3 Large Address Space Operating Systems 996

13.6.4 Networking 996

13.6.5 Parallel and Distributed Systems 997

13.6.6 Multimedia 997

13.6.7 Battery-Powered Computers 998

13.6.8 Embedded Systems 998

13.6.9 Sensor Nodes 999

13.7 SUMMARY 999

14 READING LIST AND BIBLIOGRAPHY

1003

14.1 SUGGESTIONS FOR FURTHER READING 1003

14.1.1 Introduction and General Works 1004

14.1.2 Processes and Threads 1004

14.1.3 Memory Management 1005

14.1.4 Input/Output 1005

14.1.5 File Systems 1006

14.1.6 Deadlocks 1006

14.1.7 Multimedia Operating Systems 1006

14.1.8 Multiple Processor Systems 1007

14.1.9 Security 1008

14.1.10 Linux 1010

14.1.11 Windows Vista 1010

14.1.12 The Symbian OS 1011

14.1.13 Design Principles 1011

14.2 ALPHABETICAL BIBLIOGRAPHY 1012

INDEX

1045